

```
In [1]: import numpy as np
import pandas as pd
import json
import sklearn
import eli5
import seaborn as sns
import matplotlib.patches as mpatches
import matplotlib.pyplot as plt

from sklearn.pipeline import make_pipeline
from lime.lime_text import LimeTextExplainer
from sklearn.linear_model import LogisticRegression
from sklearn.metrics import confusion_matrix
from sklearn.feature_extraction.text import CountVectorizer, TfidfVectorizer
from sklearn.model_selection import train_test_split

import warnings
warnings.filterwarnings("ignore")
```

```

In [12]: def visualize_gender_references_frequency():
    ''' Returns matplotlib plot summarizing frequency of gender mentions '''

    # set the figure size
    plt.figure(figsize=(12, 6))

    # Extract percentage
    percentage_male_female = (summary_gender.pct_male + summary_gender.pct_female) / 100
    df_perc = pd.DataFrame({'stars': ['1 star', '2 stars', '3 stars', '4 stars', '5 stars'],
                           'Gender references': round((percentage_male_female * 100, 3))})
    df_total = pd.DataFrame({'stars': ['1 star', '2 stars', '3 stars', '4 stars', '5 stars'],
                           'Gender references': [100, 100, 100, 100, 100]})

    # Create 2 bars
    bar1 = sns.barplot(x="stars", y="Gender references", data=df_total, color='Gray')
    bar2 = sns.barplot(x="stars", y="Gender references", data=df_perc, color='Black')

    # add legend
    top_bar = mpatches.Patch(color='Gray', label='Gender reference = No')
    bottom_bar = mpatches.Patch(color='Black', label='Gender reference = Yes')
    plt.legend(handles=[top_bar, bottom_bar])

    # show the graph
    return plt.show()

def visualize_gender_references_split():
    ''' Returns matplotlib plot summarizing split of genders '''

    # set the figure size
    plt.figure(figsize=(12, 6))

    # Extract percentage
    percentage_male = summary_gender.pct_male / (summary_gender.pct_male + summary_gender.pct_female)
    df_male = pd.DataFrame({'stars': ['1 star', '2 stars', '3 stars', '4 stars', '5 stars'],
                           'Gender': round((percentage_male * 100, 3))})
    df_total = pd.DataFrame({'stars': ['1 star', '2 stars', '3 stars', '4 stars', '5 stars'],
                           'Gender': [100, 100, 100, 100, 100]})

    # Create 2 bars
    bar1 = sns.barplot(x="stars", y="Gender", data=df_total, color='Pink')
    bar2 = sns.barplot(x="stars", y="Gender", data=df_male, color='Blue')

```

```

)

# add legend
top_bar = mpatches.Patch(color='Pink', label='Female')
bottom_bar = mpatches.Patch(color='Blue', label='Male')
plt.legend(handles=[top_bar, bottom_bar])

# show the graph
return plt.show()

def run_logistic_reg(df):
    '''Takes in pd df, then creates embeddings, trains logistic regression and returns most important parameters'''

    # Split train and test set
    training_data, test_data = train_test_split(df, train_size=0.8, random_state=123)
    te_y = test_data['stars']

    # Creating unigram + bigram embeddings
    vectorizer = TfidfVectorizer(stop_words='english', ngram_range=(1, 2),
                                min_df=3, lowercase=True, max_features=100000)
    bow_representation = vectorizer.fit_transform(training_data['clean_text'])
    bow_representation_test = vectorizer.transform(test_data['clean_text'])

    best_logit = LogisticRegression(C=1, solver='liblinear',
                                    penalty='l1', max_iter=1000).fit(bow_representation, training_data['stars'])

    # predict
    y_test_pred = best_logit.predict(bow_representation_test)

    # Evaluate model
    c_matrix_test = confusion_matrix(te_y, y_test_pred)
    # Accuracy
    acc = np.round(sklearn.metrics.accuracy_score(te_y, y_test_pred), 5)
    # Precision
    prec = np.round(sklearn.metrics.precision_score(te_y, y_test_pred, average=None), 3)
    prec_micro = np.round(sklearn.metrics.precision_score(te_y, y_test_pred, average='micro'), 5)
    # Recall
    rec = np.round(sklearn.metrics.recall_score(te_y, y_test_pred, average=None), 3)
    rec_micro = np.round(sklearn.metrics.recall_score(te_y, y_test_pred, average='micro'), 5)
    # F1
    f1 = np.round(sklearn.metrics.f1_score(te_y, y_test_pred, average=None), 3)
    f1_micro = np.round(sklearn.metrics.f1_score(te_y, y_test_pred, average='micro'), 5)

```

```

# Print model results
print('Acc: ', acc, ' Prec: ', prec, ' Rec: ', rec, ' f1: ', f1)

# Extract vector names
# feature_names = vectorizer.get_feature_names_out()
feature_names = vectorizer.get_feature_names()

# Create summary pd
top_x, top_y = 50, 50
weights = eli5.show_weights(estimator=best_logit, top=(top_x, top_y
),
                                target_names=training_data['stars'])
result = pd.read_html(weights.data)[0]
result = result.drop([top_x, (top_x + 1)], axis=0)
result['feature_number'] = list(map(lambda x: int(x[1:]), result.Fea
ture))
result['feature_name'] = list(map(lambda x: feature_names[x], result
.feature_number))
result['weight_num'] = list(
    map(lambda x: np.where(x[0] == "+", float(x[1:]), float(x[1:]) *
-1), result['Weight?']))

return result

def visualize_feature_importance(df, review_index):
    ''' Returns word by word importance for one given review'''
    vectorizer = TfidfVectorizer(stop_words='english', ngram_range=(1, 2
),
                                min_df=3, lowercase=True, max_features=
100000)
    bow_representation = vectorizer.fit_transform(df['clean_text'])
    bow_representation_test = vectorizer.transform(df['clean_text'])

    best_logit = LogisticRegression(C=1, solver='liblinear',
                                    penalty='l1', max_iter=1000).fit(bow
_representation,
                                                                    df[
'stars'])

    class_names = {1: '1_star', 5: '5_star'}
    LIME_explainer = LimeTextExplainer(class_names=class_names)
    c = make_pipeline(vectorizer, best_logit)

    LIME_exp = LIME_explainer.explain_instance(female_df.text[review_ind
ex], c.predict_proba)
    # print results
    print('Document id: %d' % review_index)
    print('Review: ', female_df.text[review_index])
    print('Probability 5 star =', c.predict_proba([female_df.text[review
_index]]).round(3)[0, 1])
    print('True class: %s' % class_names.get(female_df.stars[review_inde
x]))

    return LIME_exp.show_in_notebook(text=True)

```

```
In [3]: # Read in reviews
reviews = []
with open('/Users/louisgenereux/Desktop/Term 4/Text_analytics/yelp_data/et/' \
          'yelp_academic_dataset_review.json') as json_file:
    for rec in json_file:
        dic = json.loads(rec)
        reviews.append(dic)
print("- JSON format review data has been read")
```

- JSON format review data has been read

```
In [4]: # Convert to pd
reviews_df = pd.DataFrame.from_records(reviews)
print("- Data converted to pd format")
```

- Data converted to pd format

```
In [5]: # Produce summary of df
label = reviews_df['stars'].value_counts().index
observed = reviews_df['stars'].value_counts()
pct = reviews_df['stars'].value_counts() / len(reviews_df['stars'])
summary = pd.DataFrame({'stars': label,
                        'observed': observed,
                        'pct': round(pct, 4) * 100})
summary = summary.sort_values(['stars'])
```

```
In [6]: # Reading in pre-processed gendered df
gender_mentionned = pd.read_csv('yelp_gendered.csv')
gender_mentionned_unique = gender_mentionned[(gender_mentionned['male_present'] +
                                                gender_mentionned['female_present']) == 1]
print("- Gendered CSV read, entries referencing both genders are removed")
```



- Gendered CSV read, entries referencing both genders are removed

```
In [7]: # Producing summary
summary_gender = gender_mentionned_unique.groupby(['stars']).agg({
    'male_present': [ ('sum') ],
    'female_present': [ ('sum'
    ) ] })
summary_gender['pct_male'] = round(100*summary_gender['male_present']['sum']/
    list(summary['observed'].values),3)
summary_gender['pct_female'] = round(100*summary_gender['female_present']
    [ 'sum' ] /
    list(summary['observed'].values),3
)
print("- Summary of corpus: ")
print(summary_gender)
```

```
- Summary of corpus:
      male_present  female_present  pct_male  pct_female
stars
1.0           96876           100370     7.672     7.948
2.0           39107           40362     5.497     5.674
3.0           37504           34401     4.047     3.712
4.0           68114           52955     3.548     2.758
5.0          144697           93389     3.793     2.448
```

```
In [8]: # Produce visuals of dataset stats
visualize_gender_references_frequency()
visualize_gender_references_split()

# Create male and female only corpa
male_df = gender_mentionned_unique[(gender_mentionned_unique['male_prese
nt'] == 1) &
                                   (gender_mentionned_unique['stars'].isin([1, 5]))]
female_df = gender_mentionned_unique[(gender_mentionned_unique['female_p
resent'] == 1) &
                                       (gender_mentionned_unique['stars'].isin([1, 5]))]
print("- Male only and female only gendered corpa created")
```



- Male only and female only gendered corpa created

```
In [9]: # Identify top word predictors of star ratings in gendered corpus
print("- Running LOGISTIC REGRESSION")
print("- Testing result for men:")
result_men = run_logistic_reg(male_df)
print("- Testing result for women:")
result_women = run_logistic_reg(female_df)
print('- Top predictors of high ratings male')
print(result_men[0:10])
print('- Top predictors of low ratings female')
print(result_women[90:100])
```

- Running LOGISTIC REGRESSION

- Testing result for men:

Acc: 0.96608 Prec: [0.959 0.971] Rec: [0.956 0.973] f1: [0.958 0.972]

- Testing result for women:

Acc: 0.96754 Prec: [0.969 0.966] Rec: [0.969 0.966] f1: [0.969 0.966]

- Top predictors of high ratings male

	Weight?	Feature	feature_number	feature_name	weight_num
0	+29.459	x1514	1514	amazing	29.459
1	+28.732	x20195	20195	delicious	28.732
2	+23.495	x35392	35392	great	23.495
3	+22.874	x26315	26315	excellent	22.874
4	+22.726	x4562	4562	awesome	22.726
5	+22.596	x6391	6391	best	22.596
6	+21.920	x63281	63281	perfect	21.92
7	+20.010	x27827	27827	fantastic	20.01
8	+17.653	x63825	63825	phenomenal	17.653
9	+17.309	x39551	39551	highly recommend	17.309

- Top predictors of low ratings female

	Weight?	Feature	feature_number	feature_name	weight_num
92	-15.547	x7496	7496	bland	-15.547
93	-17.150	x92519	92519	used love	-17.15
94	-17.445	x21951	21951	disgusting	-17.445
95	-19.377	x4977	4977	awful	-19.377
96	-19.538	x86084	86084	terrible	-19.538
97	-21.051	x21837	21837	disappointing	-21.051
98	-21.170	x39454	39454	horrible	-21.17
99	-21.360	x92099	92099	unprofessional	-21.36
100	-28.067	x73041	73041	rude	-28.067
101	-33.179	x98796	98796	worst	-33.179


```

In [10]: # Extracting positive and negative words in the male and female situation
n
results_pos_men = list(result_men[0:50].feature_name.values)
results_pos_women = list(result_women[0:50].feature_name.values)
results_neg_men = list(result_men[51:100].feature_name.values)
results_neg_women = list(result_women[51:100].feature_name.values)

# Identifying which words appear only in subset
men_only_pos = list(set(results_pos_men) - set(results_pos_women))
women_only_pos = list(set(results_pos_women) - set(results_pos_men))
men_only_neg = list(set(results_neg_men) - set(results_neg_women))
women_only_neg = list(set(results_neg_women) - set(results_neg_men))

# Print unique words
print('- Words that are POSITIVE predictors of strong ratings when MEN are mentioned only:')
print(men_only_pos)
print('- Words that are POSITIVE predictors of strong ratings when WOMEN are mentioned only:')
print(women_only_pos)
print('- Words that are NEGATIVE predictors of strong ratings when MEN are mentioned only:')
print(men_only_neg)
print('- Words that are NEGATIVE predictors of strong ratings when WOMEN are mentioned only:')
print(women_only_neg)

```

```

- Words that are POSITIVE predictors of strong ratings when MEN are mentioned only:
['satisfied', 'notch', 'did charge', 'good', 'nicest', 'happier', 'hesitate', 'hesitation', 'exceptional', 'saved', 'pleasantly', 'enjoyed']
- Words that are POSITIVE predictors of strong ratings when WOMEN are mentioned only:
['delightful', 'tasty', 'hooked', 'exactly', 'thankful', 'sweetest', 'loved', 'beautiful', 'accommodating', 'efficient', 'fabulous', 'gorgeous']
- Words that are NEGATIVE predictors of strong ratings when MEN are mentioned only:
['condescending', 'wtf', 'excuse', 'shame', 'incompetent', 'tasteless', 'negative stars', 'response', 'ignored', 'wasted', 'unfriendly', 'flavorless']
- Words that are NEGATIVE predictors of strong ratings when WOMEN are mentioned only:
['lack', 'ridiculous', 'poisoning', 'unhelpful', 'wanted love', 'charged', 'attitude', 'ruined', 'left', 'asked', 'zero', 'acted']

```

```
In [13]: # Which data entries have the word attitude
female_df['attitude'] = list(map(lambda x: np.where(('attitude' in x), 1
, 0), female_df.text))

# Visualize feature importance
visualize_feature_importance(female_df ,762)
```

Document id: 762

Review: I don't know what it is but the coffee I get just tastes so good from here. I usually go for large coffee with almond milk.

However, the 5 starts are for the morning crew that are there around 8:30Am. You guys are absolutely wonderful. Especially the lady with the blue shirt on (manager maybe) she's just the sweetest and happiest person ever. It's not fake happy either. I was a Starbucks fan before but I've just turned to team Dunkin!!!

Looking forward to my next coffee and thank you for the wonderful service!!

Probability 5 star = 0.998


True class: 5_star

Prediction probabilities



1

5



Text with highlighted words

I don't know what it is but the coffee I get just tastes so good from here. I usually go for large coffee with almond milk.

However, the 5 starts are for the morning crew that are there around 8:30Am. You guys are absolutely wonderful. Especially the lady with the blue shirt on (manager maybe) she's just the sweetest and happiest person ever. It's not fake happy either. I was a Starbucks fan before but I've just turned to team Dunkin!!!

Looking forward to my next coffee and thank you for the wonderful service!!